
BIFROST: AGENT-BASED PUBLIC TRANSPORTATION OPTIMIZATION

Leonardo Artiles Montero
Universidad de la Habana
La Habana
leonardo16am@gmail.com

Edian Broche Castro
Universidad de la Habana
La Habana
edianbc@gmail.com

ABSTRACT

This paper presents Bifrost, a simulation project aimed at optimizing public transportation in a city using a multi-agent approach. Bifrost simulates agents representing travelers, modeled with the Belief-Desire-Intention (BDI) architecture, and utilizes Particle Swarm Optimization (PSO) to minimize Conditional Value at Risk at the 90th percentile (CVaR₉₀) of travel times. The simulation is designed to assess and improve the efficiency of public transport systems by modeling real-world behaviors and optimizing the transportation network.

1 Introduction

Efficient public transportation is essential for reducing congestion, minimizing emissions, and improving urban quality of life. Bifrost is a simulation-based approach to evaluate and optimize public transport efficiency by simulating agents in a city using the BDI framework. The agents represent travelers making transportation decisions in an environment consisting of various modes of transit. The optimization goal focuses on reducing the CVaR₉₀ of agent travel times using PSO, which helps to improve both average travel times and reliability. For experimental purposes, we used the map of Havana, Cuba.

2 Problem Statement

The problem involves optimizing public transportation in a city represented as a map, with a set of routes serviced by buses and individuals who need to move throughout the city. The challenge is to find an optimal set of bus routes that minimizes a cost function, specifically focusing on minimizing travel times while ensuring reliability. The reliability is measured by the CVaR₉₀ metric, which reflects the worst-case travel scenarios. This requires understanding individual behaviors and optimizing the routes to achieve an efficient and dependable transportation network.

3 Problem Modelation

3.1 Graph Modelation

The city is represented as a graph, where intersections of streets are nodes and edges are the streets. In the graph containing buses, the representation is similar, but for each bus stop, a new node is created, and edges are added between bus stops with less cost than walking. Additionally, a directed edge is added from the regular nodes to the bus stop nodes to represent the approximate time to wait for a bus, and another directed edge with zero cost is added in the opposite direction to represent leaving a bus. This 'walking + bus' representation only occurs in the beliefs of the agents (further details are provided later).

3.2 Route Modelation

Routes are sets of nodes, which represent the bus stops. Each bus travels using the shortest distance between stops, and for each route, there is an assigned number of buses.

3.3 Agents Modelation

The agents in the simulation are based on the BDI architecture, which allows them to make decisions in a way that mimics human behavior.

- **Beliefs:** Our agents' beliefs are the knowledge about the buses in the city, more specifically the time they have to wait at a specific bus stop, and the time a bus takes to go from one bus stop to another. More formally, the beliefs of our agents are represented as a graph, containing the time to go from one node to another in minutes, essentially the graph of the city but with distances represented as the time it takes for that specific person to walk. Additionally, more nodes are added, one for each bus stop for each route, which is connected to the original graph node via two directed edges. One edge represents the time the person believes they have to wait, and the other is zero (representing the time needed to leave the bus). There are also weights between stops, indicating the time it takes for the bus to travel from one stop to another.
- **Desires:** The agent's desires represent their ultimate goals, such as reaching a destination in the shortest possible time, minimizing travel cost, or maximizing convenience. These desires are influenced by the agent's beliefs and personal attributes. For example, an agent with limited financial resources may prioritize minimizing cost, while an agent in good physical condition may be more inclined to walk if the distance is reasonable.
- **Intentions:** Intentions in our project are the specific actions that agents decide to take in order to achieve their desires. Based on their beliefs and using fuzzy logic to account for uncertainties, agents may intend to either walk or take a bus. For example, if an agent perceives that the bus delay is too long, they may intend to walk to their destination. Conversely, if the distance is far and the agent believes their physical state is weak, they will intend to take a bus. These intentions are influenced by the agent's attributes and the current environment, guiding the agent's behavior in the simulation.
- **Fuzzy Logic:** Given the beliefs of a person, we use fuzzy logic to decide which action to take, such as whether to take a bus or walk. Each agent has unique attributes, including:
 - **Distance:** Represents the distance to be traveled, categorized as 'Close', 'Medium', or 'Far'.
 - **Bus Delay:** The expected delay for a bus, categorized as 'Short', 'Medium', or 'Long'.
 - **Physical State:** Represents the physical fitness of the agent, categorized as 'Weak', 'Average', or 'Strong'.
 - **Money:** Represents the agent's financial state, categorized as 'Low', 'Medium', or 'High'.
 - **Decision:** Represents the final action that the agent takes, which can be 'Walk' or 'Bus'.

The fuzzy rules used in the decision-making process are as follows:

- If **Distance** is 'Far' and **Physical State** is 'Weak', then **Decision** is 'Bus'.
- If **Money** is 'Low', then **Decision** is 'Walk'.
- If **Bus Delay** is 'Long', then **Decision** is 'Walk'.
- If **Physical State** is 'Strong' and **Distance** is 'Close', then **Decision** is 'Walk'.
- If **Distance** is 'Medium' and **Physical State** is 'Average', then **Decision** is 'Walk'.
- If **Money** is 'High', then **Decision** is 'Bus'.

These fuzzy rules, defined in the fuzzy logic engine, help agents make decisions based on their attributes and beliefs about the environment. For example, if an agent perceives the distance to be far and their physical state to be weak, they are more likely to choose taking a bus. Conversely, if they perceive a long bus delay, they may opt to walk.

4 Optimization

4.1 Simulation

The simulation takes place over several days (usually 7). At the start of the simulation, agents have beliefs about the waiting times for each bus and the time buses take between stops—very basic and general knowledge. Each day, the

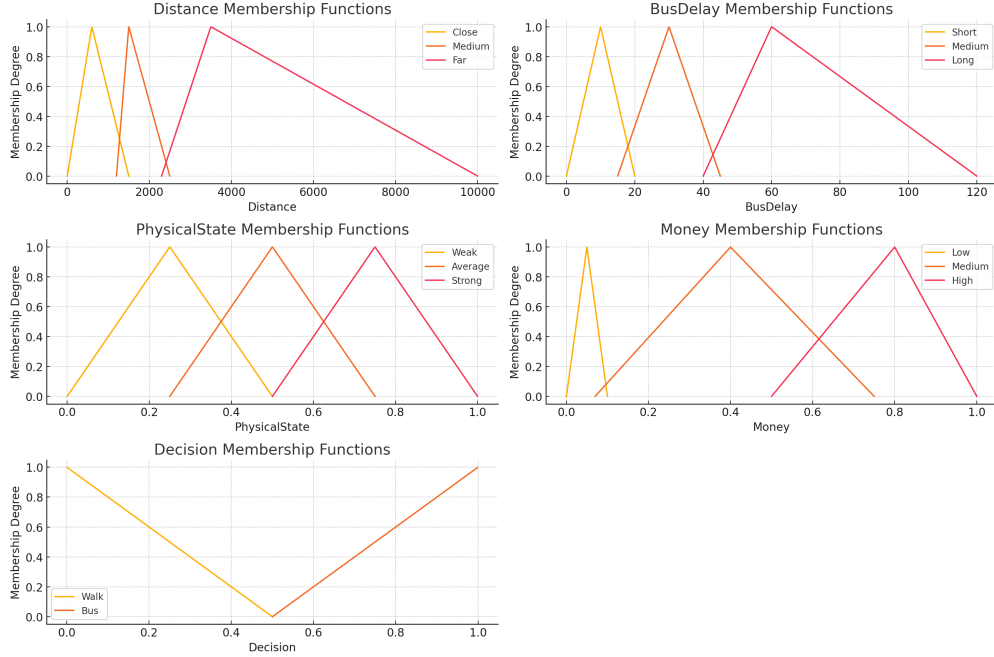


Figure 1: Fuzzy Logic

agents move around the city, and at the end of the day, based on their experience with the buses they took, they update their beliefs with what they learned.

The movement of agents through the city is an event-based simulation, where there are two basic entities: the people and the buses. The buses follow given routes and move in cycles through their respective stops. The people, on the other hand, decide where they want to go and, based on their beliefs, plan the best route and take it. They board the bus when it arrives and get off when they reach their destination stop. A person may take more than one bus if necessary. This simulation involves the following events defined in the code:

- **PERSON_START_WALKING:** An agent starts walking towards a bus stop or destination.
- **PERSON_ARRIVE_STOP:** An agent arrives at a bus stop and waits for a bus.
- **BUS_ARRIVE_STOP:** A bus arrives at a stop, allowing agents to board or disembark.
- **PERSON_GET_OFF_BUS:** An agent gets off a bus at their desired stop.
- **PERSON_ARRIVE_WORK:** An agent reaches their final destination, such as work or another point of interest.

At the end of each day, the agents record the bus delay times and the times between stops and average them with the previously known times, except for the first day, where they simply keep the newly gained knowledge (as if they had never taken that bus before).

4.2 Error Metric

The primary metric used for optimization is $CVaR_{90}$, which measures the Conditional Value at Risk at the 90th percentile. Specifically, it represents the average of the worst 10% of travel times. This metric was chosen because it not only considers average performance but also focuses on the tail end of the distribution, ensuring that the most problematic travel scenarios are addressed. By minimizing $CVaR_{90}$, we aim to provide a more reliable transportation system that performs well even in the worst-case situations, improving both efficiency and user satisfaction.

4.3 Particle Swarm Optimization (PSO)

Particle Swarm Optimization (PSO) is an optimization algorithm inspired by the social behavior of birds flocking or fish schooling. In our project, PSO is used to optimize the bus routes to minimize the $CVaR_{90}$ metric of travel times. PSO

works by initializing a swarm of particles, where each particle represents a potential solution (i.e., a set of routes). These particles move through the solution space, adjusting their positions based on their own experience and the experience of neighboring particles to find the optimal solution.

For each route in our simulation, we used 5 intermediate points to optimize the path, with the remaining stops being positioned along the shortest path between these intermediate points. This approach allows us to maintain flexibility in route design while avoiding overly simplistic solutions. In earlier experiments, we used only two intermediate points, which resulted in less effective optimization. In real-world settings, bus routes often do not follow the shortest path between two points, as they need to service multiple stops to meet demand. By using more intermediate points, we allow routes to diverge from the shortest path where necessary, leading to a more realistic and efficient public transportation system.

5 Experiments

5.1 Comparison of Results with Different PSO Hyperparameters

In this experiment, we tested the Particle Swarm Optimization (PSO) algorithm to optimize public transportation based on the CVaR90 metric. The results were compared using different hyperparameter configurations, specifically the number of particles and the PSO acceleration coefficients. Figure 2 shows the evolution of the CVaR90 error over training epochs, considering different PSO configurations.

5.1.1 Tested Configurations

The tested configurations differ in:

- **Number of particles:** Three values were tested: 5, 10, and 15 particles.
- **Acceleration coefficients (cognitive and social):** Two sets of acceleration coefficients were used:
 - (1.5, 2): This configuration places a greater emphasis on the social component.
 - (3, 4): This configuration gives more weight to both the individual (*cognitive*) and social exploration.

5.1.2 Interpretation of the Graph

Figure 2 shows how the CVaR90 error evolves during 14 training epochs for the different combinations of particles and acceleration coefficients:

- **5 particles (1.5, 2)** (blue dashed line): Starts with a CVaR90 error of approximately 115, but stabilizes quickly, with no significant improvement after the second epoch.
- **10 particles (1.5, 2)** (green dashed line): Shows a similar behavior to 5 particles, starting with an error around 120, but achieves a slight reduction, stabilizing after 6 epochs.
- **15 particles (1.5, 2)** (red dashed line): Although it starts with a lower CVaR90 (around 105), it shows faster improvement, constantly reducing until reaching an error close to 100.
- **5 particles (3, 4)** (cyan dashed line): This configuration starts at an initial error of approximately 110 and improves quickly, reaching the best performance among all configurations, with a final error close to 98.
- **10 particles (3, 4)** (green solid line): Although it starts with a CVaR90 around 115, it gradually improves to a final error close to 100.
- **15 particles (3, 4)** (purple dashed line): Begins with a CVaR90 around 105 but does not show substantial improvement throughout the epochs, suggesting that this configuration may not be ideal.

5.1.3 Conclusions

- All configurations demonstrated similar performance, likely due to the high dimensionality of the search space, which makes the optimization more challenging.
- Although the configuration with **5 particles (3, 4)** achieved the best result in terms of CVaR90 minimization, the differences between configurations were marginal.
- The results suggest that the number of particles tested may have been insufficient to fully explore the solution space, which could be why the larger configurations with 10 and 15 particles did not significantly outperform the smaller configurations. Ideally, experiments should be conducted with much larger numbers of particles to better navigate the high-dimensional search space.

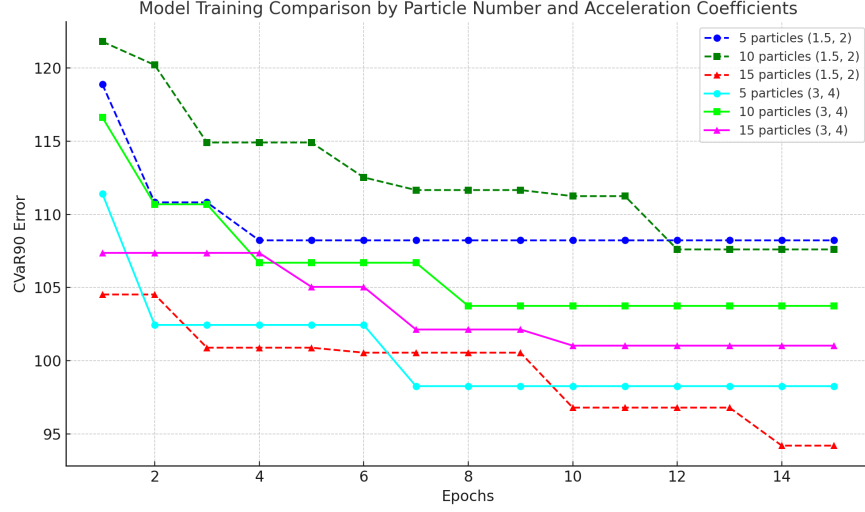


Figure 2: Model training comparison using different PSO configurations (particle numbers and acceleration coefficients).

5.2 Sanity Checks

To ensure that our model behaves as expected under extreme and simplified conditions, we conducted two sanity check experiments. These experiments are designed to validate the basic functionality of the optimization process before tackling more complex scenarios.

5.2.1 Experiment 1: Three Routes for Many People

In this experiment, we asked the question: *What would happen if there were only two bus routes available for a large number of people spread across the city?* The goal was to observe how the system would handle the optimization of such limited resources.

As shown in Figure ??, the three routes were positioned in a roughly X-shaped pattern, crossing the city and covering as much area as possible. This result is intuitive, as the model attempts to maximize coverage with the limited number of routes, ensuring that the majority of people are served by at least one of the routes.

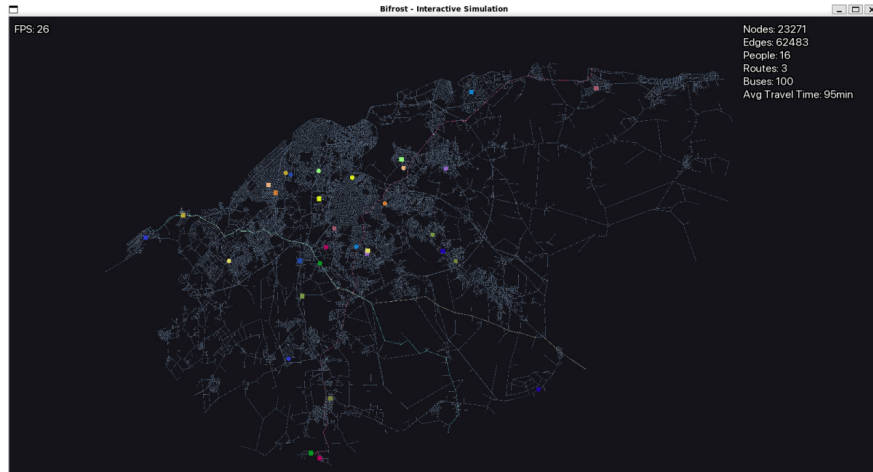


Figure 3: Route configuration for three routes serving many people.

5.2.2 Experiment 2: One Person and One Route

The second sanity check asked: *What would happen if there were only one person and one bus route in the city?* The objective was to see how the optimization would perform in this simplified scenario where the needs of a single individual dominate the solution.

As expected, the bus route was established approximately along the path between the person's home and their workplace, as illustrated in Figure 4. This outcome confirms that the model behaves logically in extreme cases by minimizing the travel time for the single individual.



Figure 4: Route configuration for one person and one route between their home and workplace.

5.2.3 Conclusions

Both experiments validate the basic operation of the model. In the case of many people and limited routes, the model optimizes coverage across the city, while in the case of a single person, the route is tailored to their specific needs. These sanity checks give us confidence that the model is functioning as intended under simplified circumstances, allowing us to proceed with more complex scenarios.

5.2.4 Experiment 3: Changing Bus Availability

In this experiment, we created a scenario with a single person and a single bus route. The route was initially equipped with three buses and roughly followed the path between the person's home and their workplace. The person's behavior was simple: they commuted daily from their home to their workplace and only used the bus route available.

Initially, the person consistently took the bus as part of their commute. However, every three days, one of the buses was removed from service. This reduction in bus availability eventually affected the person's commuting decisions.

By day 8, the person realized that walking to work was faster than waiting for the increasingly delayed bus. Specifically, on day 8, the bus was delayed to the point where it would take 43.439 minutes for the person to reach work by bus, compared to only 31.122 minutes if they walked. Having observed the pattern of delays over the previous two days, the person decided to stop using the bus and opted to walk to work from that point onward.

This experiment illustrates how agents adapt their behavior based on their experiences with transportation systems and make rational decisions to minimize travel time.

6 Further Improvements

Future improvements to Bifrost include:

- Exploring other optimization algorithms beyond PSO.
- Adding schedules to bus routes.
- Optimizing the project to simulate a greater number of agents, potentially using parallelization with GPUs.
- Using more realistic data for agents' geographical dispersion.
- Modeling the problem in a more flexible way to allow modifications to the number of buses per route during optimization.
- Minimizing additional variables, such as fuel consumption or traffic congestion.

7 General Conclusions

Bifrost demonstrates a novel approach to optimizing public transportation through agent-based simulation and Particle Swarm Optimization (PSO). By focusing on minimizing CVaR_{90} , the model addresses not only the average travel times but also the worst-case scenarios, ensuring that the transportation system is both efficient and reliable. This approach makes it possible to create a balanced network that serves the needs of different users effectively, particularly those affected by delays and less optimal conditions. Moreover, this project could have significant applications in real-world traffic optimization, offering insights into how public transit systems can be made more efficient and reliable, ultimately improving urban mobility and reducing congestion.